

CIS570 Modern Programming Language Implementation

Instructor: E Christopher Lewis
Office hours: TDB
605 Levine
eclewis@cis.upenn.edu

Admin. Assistant: Cheryl Hickey
cherylh@central.cis.upenn.edu
502 Levine

URL: <http://www.cis.upenn.edu/~eclewis/cis570>

What is CIS570 About?

Program representation

- Can we do better than ASCII .c files?

Analysis

- How do we mechanically derive meaning and intent?
- How do we reason about programs?

Transformation

- How do we use results of analysis to make the representation “better”?

Domains

- Language implementation (including **Compilation**)
- Computation understanding
 - Software engineering tools (bug detectors, *etc.*)

Structure of CIS570

Lectures

- Participation is essential (no text book)

Reading

- 6 or 7 papers
- Read and answer discussion questions before designated class

Discussions

- Some classes will be devoted to discussion
- But there is always room for discussion

Exams


- Final perhaps?

Homework

- Answer discussion questions before class

Bottom line

- Very efficient class



Plan for Today

Motivation

- Why study compilers?

Issues

- Look at some sample optimizations and assorted issues

Administrivia

- Course details

Motivation

What is a compiler?

- A translator that converts a source program into an target program

What is an optimizing compiler?

- A translator that *somehow* improves the program (versus non-opt. comp.)

Why study compilers?

- They are *specifically* important:
Compilers provide a bridge between applications and architectures
- They are *generally* important:
Compilers encapsulate techniques for reasoning about programs and their behavior
- They are cool:
First major computer application

Traditional View of Compilers

Compiling down

- Translate high-level language to machine code

High-level programming languages

- Increase programmer productivity
- Improve program maintenance
- Improve portability

Low-level architectural details

- Instruction set
- Addressing modes
- Pipelines
- Registers, cache, and the rest of the memory hierarchy
- Instruction-level parallelism

Isn't Compilation A Solved Problem?

“Optimization for scalar machines is a problem that was solved ten years ago”

-- David Kuck, 1990

Applications keep changing

- Interactive, real-time, mobile, secure

Machines keep changing

- New features present new problems (*e.g.*, MMX, EPIC, profiling support, TM)
- Changing costs lead to different concerns (*e.g.*, mem v. ALU ops)

Values keep changing

- Correctness
- Run-time performance
- Code size
- Compile-time performance
- Power
- Security

Languages keep changing

- Wacky ideas (*e.g.*, OOP and GC) have gone mainstream

Modern View of Compilers

Analysis and translation are useful everywhere

- Analysis and transformations can be performed at run time and link time, not just at “compile time”
- Translation can be used to improve security
- Analysis can be used in software engineering
 - Program understanding
 - Reverse engineering
- Increased interaction between hardware and compilers can improve performance
- Bottom line
 - Analysis and transformation play essential roles in computer systems
 - Computation important \Rightarrow *understanding* computation important

Types of Optimizations

Definition

- An *optimization* is a transformation that is expected to improve the program in some way; often consists of *analysis* and *transformation* e.g., decreasing the running time or decreasing memory requirements

Machine-independent optimizations

- Eliminate redundant computation
- Move computation to less frequently executed place
- Specialize some general purpose code
- Remove useless code

Types of Optimizations (cont)

Machine-dependent optimizations

- Replace a costly operation with a cheaper one
- Replace a sequence of operations with a cheaper one
- Hide latency
- Improve locality
- Reduce power consumption


Enabling transformations

- Expose opportunities for other optimizations
- Help structure optimizations

Sample Optimizations

Arithmetic simplification



- Constant folding

e.g., `x = 8/2;`  `x = 4;`


- Strength reduction

e.g., `x = y * 4;`  `x = y << 2;`

Constant propagation


– e.g.,
`x = 3;`
`y = 4+x;`  `x = 3;`
`y = 4+3;`  `x = 3;`
`y = 7;`

Copy propagation

– e.g.,
`x = z;`
`y = 4+x;`  `x = z;`
`y = 4+z;`

Sample Optimizations (cont)

Common subexpression elimination (CSE)


– e.g.,
`x = a + b;`
`y = a + b;`  `t = a + b;`
`x = t;`
`y = t;`

Dead (unused) assignment elimination

– e.g.,
`x = 3;`
... `x` not used...
`x = 4;` 

this assignment is dead


Dead (unreachable) code elimination

– e.g.,
`if (false == true) {`
 `printf("debugging...");`
`}` 

this statement is dead

Sample Optimizations (cont)


Loop-invariant code motion

– e.g., `for i = 1 to 10 do`
 `x = 3;`
 ...  `x = 3;`
 `for i = 1 to 10 do`
 ...

Induction variable elimination

– e.g., `for i = 1 to 10 do`  `for p = &a[1] to &a[10] do`
 `a[i] = a[i] + 1;` `*p = *p + 1`

Loop unrolling

– e.g., `for i = 1 to 10 do`  `for i = 1 to 10 by 2 do`
 `a[i] = a[i] + 1;` `a[i] = a[i] + 1;`
 `a[i+1] = a[i+1] + 1;`

Is an Optimization Worthwhile?

Criteria for evaluating optimizations

- **Safety**: does it preserve behavior?
- **Profitability**: does it actually improve the code?
- **Opportunity**: is it widely applicable?
- **Cost (compilation time)**: can it be practically performed?
- **Cost (complexity)**: can it be practically implemented?

Scope of Analysis/Optimizations

Peephole

- Consider a small window of instructions
- Usually machine specific
- Know nothing about context

Global (intraprocedural)

- Consider entire procedures
- Must consider branches, loops, merging of control flow
- Use data-flow analysis
- Make certain assumptions at procedure calls

Local

- Consider blocks of straight line code (no control flow)
- Simple to analyze
- Know nothing about context

Whole program (interprocedural)

- Consider multiple procedures
- Analysis even more complex (calls, returns)
- Hard with separate compilation

Limits of Compiler Optimizations

Fully Optimizing Compiler (FOC)

- $\text{FOC}(P) = P_{\text{opt}}$
- P_{opt} is the *smallest* program with same I/O behavior as P

Observe

- If program Q produces no output and never halts, $\text{FOC}(Q) = L$: goto L

Aha!


- We've solved the halting problem?!

Moral

- Cannot build FOC
- Can always build a better optimizing compiler (*full employment theorem* for compiler writers!)


Optimizations Don't Always Help

Common Subexpression Elimination

$x = a + b$		$t = a + b$
$y = a + b$		$x = t$
		$y = t$
$\underbrace{\hspace{1.5cm}}$		$\underbrace{\hspace{1.5cm}}$
2 adds		1 add
4 variables		5 variables

Optimizations Don't Always Help (cont)

Fusion and Contraction

$\text{for } i = 1 \text{ to } n$		$\text{for } i = 1 \text{ to } n$
$T[i] = A[i] + B[i]$		$t = A[i] + B[i]$
$\text{for } i = 1 \text{ to } n$		$C[i] = D[i] + t$
$C[i] = D[i] + T[i]$		$\underbrace{\hspace{1.5cm}}$
		t fits in a register, so no loads or stores in this loop.
		Huge win on most machines.
		Degrades performance on machines with hardware managed stream buffers.

Optimizations Don't Always Help (cont)

Backpatching

`o.foo();` } In Java, the address of `foo()` is often not known until runtime (due to dynamic class loading), so the method call requires a **table lookup**.

After the first execution of this statement, **backpatching** replaces the table lookup with a direct call to the proper function.

Q: How could this optimization ever hurt?

A: The Pentium 4 has a *trace cache*, when any instruction is modified, the entire trace cache has to be flushed.

Phase Ordering Problem

In what order should optimizations be performed?

Simple dependences

- One optimization creates opportunity for another
e.g., copy propagation and dead code elimination

Cyclic dependences

- *e.g.*, constant folding and constant propagation

Adverse interactions

- *e.g.*, common subexpression elimination and register allocation
e.g., register allocation and instruction scheduling

Engineering Issues

Building a compiler is an engineering activity

Balance multiple goals

- Benefit for *typical* programs
- Complexity of implementation
- Compilation speed

Overall Goal

- Identify a small set of general analyses and optimization
- Easier said than done: just one more...

Beyond Optimization

Security and Correctness

- Can we check whether pointers and addresses are valid?
- Can we detect when untrusted code accesses a sensitive part of a system?
- Can we detect whether locks are used properly?
- Can we use compilers to certify that code is “correct”?
- Can we use compilers to obfuscate code?

Next Time

Reading

- “Binary Translation” by Sites *et al.*

Lecture

- Undergraduate compilers in a day!